

Cloud Computing

Winter Term 2018/2019

Tutorial Session 2



Anton Gulenko
Complex and Distributed IT-Systems

anton.gulenko@tu-berlin.de

Practical Assignment 2

- Due: 20.12.2018
- Summary:
 - Work on 1 host machine
 - Preferably your laptop (physical machine). If you don't have Linux, use a VM.
 - Prepare virtualization environments:
 - Qemu/KVM
 - Docker
 - Write 2 new benchmarks:
 - Forksum
 - Nginx
 - Execute benchmarks on different virtualization platforms

Practical Assignment 2

- Public cloud platforms are not mandatory for this assignment, but do not yet delete your accounts (shut down your VMs if not used for the assignment)
- We will continue using OpenSubmit, **BUT**:
 - Only for validating that all assignment files are there
 - If your submission fails the test, we will still grade it, but very likely your points will be reduced

Virtualization Platforms

- Docker
 - Write one Dockerfile for every benchmark
 - The containers should contain all material for the respective benchmark, execute the benchmark when started without parameters, and exit after printing the results
 - Exception: Nginx benchmark, since it is executed outside the container
- Qemu (with and without KVM)
 - Use a cloud image of any Debian based Linux distribution (e.g. Ubuntu)
 - Load all benchmark material in one image, it will not be submitted

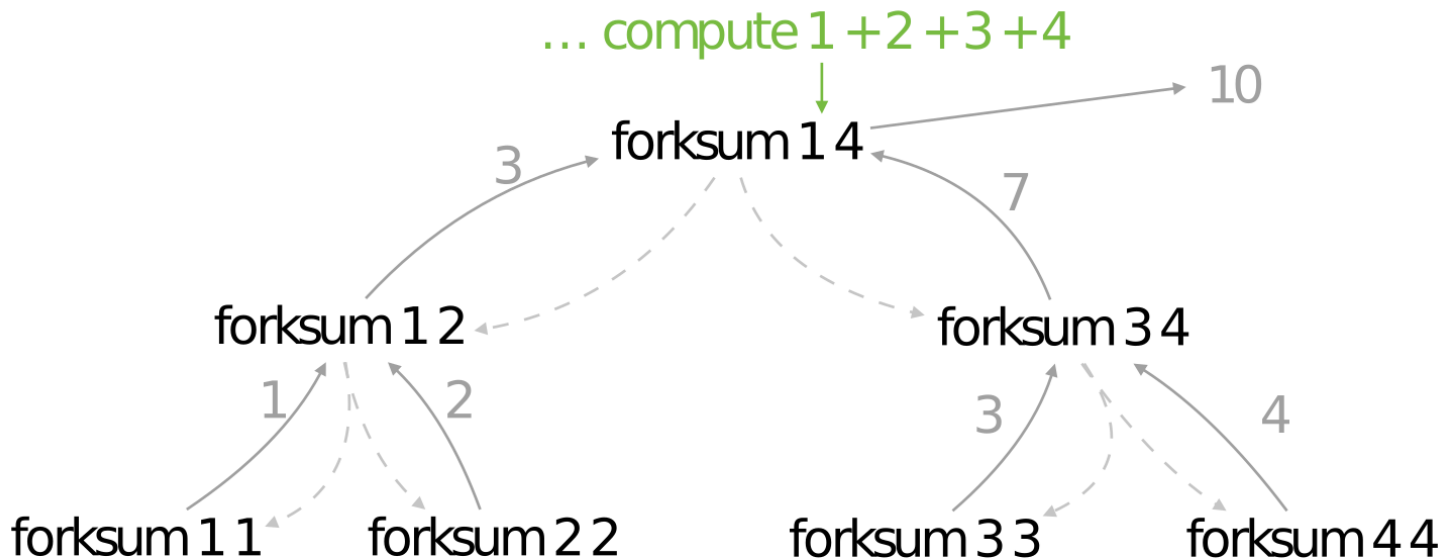
Benchmarks

- Basic resources: CPU, RAM and disk access
 - Reuse from assignment 1
 - Exception: For disk benchmark, only use fio to measure write accesses per second
- New benchmark: forksum
 - Simulates creation of many parallel processes
 - Main benchmark target: system calls
- New benchmark: Nginx web server
 - Download big file over the network
 - Main benchmark target: network performance

Forksum

- Program received 2 parameters: start and end of integer range
 - Task: compute sum of all integer within the range
 - Example: `./forksum 100 1000` should print **495550**
- Every sum is executed by a separate child process
 - If `start == end`: output value
 - Else: spawn 2 child processes: one for lower sub-range and one for upper sub-range
 - After child-processes return their results, parse them and output the sum

Forksum: Example



- - - - -> fork
—————> stdout

Forksum: required C system calls

- **fork()**: continue program as two separate processes
 - Return value of `fork()` tells the program if it is the child or parent process
- **pipe()**: create a pipe that can be used to read the standard output of a child process
- **read()**: can be used to read from a pipe
- **dup()/dup2()**: Used to assign the standard output of a child to a pipe
- **wait()**: Wait for child processes to exit
- Other useful functions: `strtol()`, `printf()`, `fprintf()`, `close()`, `exit()`,
- Example for how to read output from a child process:
 - http://www.microhowto.info/howto/capture_the_output_of_a_child_process_in_c.html
- Hint: the **exec1()** system call is not necessary

Nginx Benchmark

- Install Nginx web server
- Configure Nginx to serve static files from the disk
- Add a file (> 500MB) to the Nginx server
- Your benchmark script should receive an IP address as parameter and generate requests to the file (at least 2 requests at the same time)
- After ~10 seconds, the average access duration should be printed
- You can write a benchmarking client in any programming language or use any existing HTTP load generation tool